

FULL SEO AUDIT REPORT

kautilya-pe.com

Prepared by Claude SEO

80/100

SEO Health Score

August 18, 2026 Full Audit

Table of Contents

1. Executive Summary

Key Metrics, Critical Issues & Quick Wins

2. Audit Categories

What Works

Findings by Severity

3. Action Plan

Phased Roadmap

4. Data Sources & Methodology

1. Executive Summary

This report presents a comprehensive SEO audit of **kautilya-pe.com** for a **B2B buy-side M&A / private equity advisory firm (Mumbai, India; active across India, UK, UAE, MENA, Europe, APAC)** site, generated on August 18, 2026. Findings combine technical, content, schema, performance, visual, and search-readiness evidence as available.

80/100

SEO HEALTH SCORE

Critical Issues Found:

1. **Info:** CRITICAL: Homepage has a stuck 'Loading Sky...' WebGL asset creating ~17,500px of dead scroll space between the hero and main content (components/ThreeScene.tsx)
2. **Info:** CRITICAL: ~94MB of unoptimized images sitewide (54 sitemap images), 16 exceeding 1MB, ten Stories-page PNGs at 7-9.2MB each; zero next/image, WebP/AVIF, or responsive srcset usage
3. **Info:** CRITICAL: Deprecated HowTo schema (removed by Google Sept 2023) still live on /approach and /stories/[slug] pages, polluting Search Console
4. **Info:** HIGH: Apex domain kautilya-pe.com redirects to www with a temporary 307 instead of a permanent 301/308
5. **Info:** HIGH: news-sitemap.xml misapplies the Google News protocol to evergreen /stories/* content, several dated 2024, likely generating persistent GSC errors

Quick Wins:

1. Set apex-to-www redirect to permanent 301/308 in Vercel domain settings
2. Delete deprecated HowTo JSON-LD blocks from /approach and /stories/[slug] pages
3. Fix llms.txt FAQ count mismatch (23 -> 26 questions, one-line edit)
4. Remove news-sitemap.xml route and its robots.txt reference (stories already covered by sitemap.xml)
5. Add explicit width/height to portfolio-logo tags to close a CLS risk

2. Audit Categories

2.1 Technical SEO

Score: 87/100

What Works

- All 35 sitemap URLs return HTTP 200, no 404s or soft-404s
- Every sampled page is SSG (x-nextjs-prerender: 1, is_spa: false) with no CSR/SPA dependency
- Complete, consistent security headers sitewide (HSTS, X-Frame-Options, X-Content-Type-Options, Referrer-Policy, Permissions-Policy) via next.config.ts
- Canonical tags and meta robots are clean, unique, and correct on every page sampled
- Valid multi-type JSON-LD on every sampled page
- llms.txt and llms-full.txt both resolve and are comprehensive
- HTTP to HTTPS redirects are clean single-hop 308s

Findings

Apex domain uses a temporary 307 redirect instead of permanent 301/308 **High**

`http://kautilya-pe.com/ -> 308 -> https://kautilya-pe.com/ -> 307 -> https://www.kautilya-pe.com/` is a 3-hop chain with a non-permanent final hop, inconsistent with the clean 308 seen on the www variant.

Recommendation: Set the redirect type for `kautilya-pe.com -> www.kautilya-pe.com` to Permanent (308) in Vercel Domains settings, or add an explicit `vercel.json/next.config.ts` redirect with `permanent: true`.

No Content-Security-Policy header **Medium**

HSTS, X-Frame-Options, X-Content-Type-Options, Referrer-Policy, Permissions-Policy are all present sitewide but no CSP or CSP-Report-Only header is defined anywhere in `next.config.ts headers()`.

Recommendation: Add a CSP header scoped to known third-party origins (Google Tag Manager, the five badge domains, Unsplash) in `next.config.ts headers()`.

No IndexNow protocol support **Medium**

No IndexNow key file or submission integration found in the codebase, despite regular `/newsletter` and `/stories` publishing that would benefit from instant Bing/Yandex crawl notification.

Recommendation: Generate an IndexNow key, host it at `{key}.txt`, and add a submission call triggered on publish for `/blog/[slug]`, `/newsletter/[slug]`, `/stories/[slug]`.

Sitewide WebGL SplashCursor effect is a latent INP risk **Medium**

A continuous per-pointer-move WebGL fluid-dye simulation runs on every route via SplashCursorWrapper in the root layout. Deferred via next/dynamic so it doesn't block LCP, but competes with click/tap handlers on the main thread.

Recommendation: Lab-test INP with the cursor effect active vs. disabled; gate it to pointer-fine devices only via `window.matchMedia('(pointer: fine)')` so mobile visitors skip it entirely.

Homepage canonical omits trailing slash **Low**

Homepage is served at <https://www.kautilya-pe.com/> but its canonical tag reads <https://www.kautilya-pe.com> (no trailing slash), inconsistent with every other page's exact-match canonical.

Recommendation: Add the trailing slash to the homepage canonical to match observed serving behavior exactly.

2.2 Content Quality

Score: 82/100

What Works

- Specific, falsifiable case-study numbers (99.83% match rate, 30,134 transactions, named deal values) verified present in rendered HTML, not just marketing copy
- Full 8-person named team roster with roles and LinkedIn links on /team
- Dated, bylined long-form content with real publish dates matching JSON-LD
- Newsletter deal-teardowns are the strongest content asset (sourced stats, scored rubric, tables)
- Zero AI-pattern filler language detected on any sampled page
- Good topical depth on long-form pages (1,700-3,300 words)

Findings

Single-author concentration across all long-form content **Medium**

All 8 sampled blog/newsletter/story posts are bylined 'Dev Shah' only, despite 6 other named specialists on /team whose expertise is never attributed to specific published content.

Recommendation: Attribute posts to the team member with domain expertise, or add a 'Reviewed by' credit line, to diversify the expertise signal.

No visible per-post author bio or credentials block Medium

Articles show 'Dev Shah - [date]' only with no inline bio, headshot, or credentials snippet; the founder bio only exists on /team and in JSON-LD.

Recommendation: Add a short author byline block (1-2 sentences + photo + link to /team) at the top or bottom of every long-form post.

No visible phone number or physical address in body copy Medium

Across all 15 pages sampled, the only visible contact signal is an email address; no street address or phone number appears in rendered text anywhere.

Recommendation: Add a visible registered office address and contact phone number to /team, /engage, and the footer to strengthen Trustworthiness E-E-A-T.

Homepage body copy is thin relative to a 500-word floor Low

Stripped of nav/footer boilerplate, the homepage yields ~357 words of substantive body copy, below the typical homepage topical-coverage floor.

Recommendation: Add 2-3 more substantive paragraphs surfacing the track-record stats on-page, not just in llms.txt.

Stories case-study stats remain prose-only with no tables/visuals Low

Flagship stats-heavy case studies (e.g. /stories/inspire3) deliver numbers as prose only, with no supporting chart or table, unlike the newsletter's table pattern.

Recommendation: Add at minimum one summary-stats visual per Stories case study, matching the newsletter's table pattern.

2.3 On-Page SEO

Score: 82/100

What Works

- Consistent 'PageName | Kautilya | Keywords' title pattern across all pages sampled
- Exactly one H1 per page on every page sampled, no missing/duplicate H1s
- Every page has a unique meta description
- Canonical tags present and correct on every page, all pointing to the www host
- robots meta consistently index,follow with no accidental noindex

- OG:title distinct from and complementary to the SEO title on long-form pages

Findings

Multiple meta descriptions exceed the ~155-160 character SERP display limit **Medium**

9 of 16 sampled pages have over-length meta descriptions; worst offender is the Torrent-JB Chemicals newsletter post at 255 characters, risking truncation in search results.

Recommendation: Trim to 150-155 characters with the most compelling/keyword-relevant clause first.

/portfolio and /newsletter have zero H2 subheadings **Medium**

Both list-style pages rely on a single H1 with no subheading structure despite listing multiple distinct case studies/newsletter issues.

Recommendation: Add H2s per portfolio company / per newsletter issue for on-page structure and accessibility.

Portfolio-logo and OG images missing width/height attributes **Medium**

Portfolio company logos render as plain tags with no explicit dimensions, creating CLS risk.

Recommendation: Add explicit width/height (or migrate to next/image, which enforces this).

Third-party badge images hotlinked on key pages **Low**

Homepage and /portfolio embed external badges from five low-authority directory sites via direct hotlinked tags, adding external DNS/connection overhead.

Recommendation: Self-host or remove these badges given limited relevance to a B2B M&A advisory audience.

2.4 Schema & Structured Data

Score: 78/100

What Works

- Fully SSR'd, valid JSON on every block sampled, no malformed/truncated data
- Unusually complete Organization schema (foundingDate, areaServed, knowsAbout, founder, sameAs)
- BreadcrumbList present on every deep page sampled
- Person schema for all 8 team members with jobTitle and sameAs
- BlogPosting correctly used for blog posts with all required and recommended properties
- Absolute URLs and ISO 8601 dates throughout, correct SearchAction formatting

Findings

Deprecated HowTo schema still in production **Critical**

Google removed HowTo rich results in September 2023. This markup is live on /approach and conditionally on /stories/[slug] pages, earning nothing and polluting Search Console's Enhancements report.

Recommendation: Remove the howToSchema block from /approach and the howToLd block from stories pages entirely; leave visible HTML content in place.

FAQPage markup has no Google SERP benefit as of the audit date **Info**

Google retired FAQ rich results for all sites on 2026-05-07. FAQPage markup site-wide produces no SERP feature currently, though it may still aid AI/GEO extraction.

Recommendation: No urgent action; if kept, label internally as GEO-only, unconfirmed benefit rather than rich-result schema.

FAQ question-count mismatch between llms.txt and live FAQPage schema **Low**

llms.txt states 23 answered questions; live FAQPage JSON-LD contains 26 Question items.

Recommendation: Update llms.txt to state 26 answered questions.

Two separate top-level entities for the same business with no @id linkage **Medium**

Homepage emits both an Organization and a ProfessionalService for Kautilya independently, with drifting offer-catalog details (one includes pricing, the other doesn't) and no @id to declare they're the same entity.

Recommendation: Give the canonical Organization a stable @id (e.g. <https://www.kautilya-pe.com/#organization>) and reference it from WebSite.publisher and ProfessionalService; reconcile the pricing discrepancy.

NewsArticle type used for evergreen Stories content **Low**

Stories and Newsletter pages both use @type: NewsArticle, but Stories are evergreen case studies, not time-sensitive journalism.

Recommendation: Change Stories pages to plain Article or BlogPosting; keep NewsArticle on Newsletter posts, which do track live deal news.

Malformed empty-string servicePhone property **Low**

app/approach/page.tsx sets servicePhone: "" inside ServiceChannel, an invalid placeholder value.

Recommendation: Remove the servicePhone key entirely or populate with a real number.

2.5 Performance

Score: 85/100

What Works

- Excellent TTFB across all pages tested (55-109ms)
- Brotli compression confirmed sitewide
- Static assets cached with immutable, max-age=31536000
- Self-hosted next/font woff2 fonts with font-display: swap, no FOIT, no third-party Google Fonts round trip
- Content fully server-rendered with no CSR dependency for LCP-relevant content
- Blog images correctly use next/image with responsive srcSet and explicit dimensions

Findings

Portfolio logo images missing width/height (CLS risk) **Medium**

Two logo images on /portfolio are raw tags with no dimensions and no next/image optimization, unlike the blog template which correctly uses next/image.

Recommendation: Convert to next/image or add explicit width/height attributes.

Five third-party badge images from five external domains on every page **Low**

Badge images from five separate external origins load identically on home, /approach, /portfolio, blog, and stories pages, adding DNS/TLS overhead for content irrelevant to the target audience.

Recommendation: Self-host or lazy-load these badge images; consider removing entirely.

HTML responses use Cache-Control: max-age=0, must-revalidate **Low**

Browser is told to revalidate HTML on every navigation rather than serving from disk cache, though Vercel's ISR edge caching (300s revalidate) already compensates well given the low TTFBs measured.

Recommendation: Consider explicit s-maxage/stale-while-revalidate tuning if perceived back/forward navigation speed becomes a priority; low impact at current TTFB levels.

No CrUX/PSI field or full Lighthouse data obtained **Info**

No Google API credentials configured in this environment; PageSpeed Insights returned a rate-limit error on every attempt. All performance findings are lab-based estimates from curl timing and header inspection.

Recommendation: Re-run PageSpeed Insights / CrUX Vis once API access is available to confirm lab signals hold under real user conditions, particularly mobile INP.

2.6 Images

Score: **38/100**

What Works

- 100% alt-text coverage across all ~92 images sampled on 16 pages, zero missing or empty alt attributes
- loading=lazy used on below-the-fold avatar images
- Logo/icon assets are SVG, small footprint
- Descriptive, semantic image filenames

Findings

Story/blog images are massively oversized **Critical**

Total payload of ~94MB across the 54 images referenced in image-sitemap.xml; 16 of 54 (30%) exceed 1MB, with ten Stories-page PNGs at 7-9.2MB each (e.g. /images/stories/borderless/exhausting-market.png at 9.22MB).

Recommendation: Re-export chart/diagram images at a sane display resolution (≤ 1600 px wide), compress with a tool like squoosh/sharp, convert to WebP or AVIF, targeting <150-250KB per content image.

No modern image format usage, zero next/image optimization pipeline **High**

Every image is served as raw .png/.jpg/.jpeg directly from static paths with zero requests through /_next/image, confirming plain tags are used instead of Next.js's built-in image optimization.

Recommendation: Migrate all content/avatar/portfolio-logo usages to next/image for automatic responsive srcset, lazy-loading, CLS prevention, and WebP/AVIF serving.

No responsive srcset, identical image served to mobile and desktop **High**

srcset appears 0 times across homepage, /portfolio, and /stories HTML; a mobile visitor downloads the identical multi-megabyte asset as a desktop visitor.

Recommendation: Same fix as next/image migration above.

Missing width/height on portfolio-logo and badge images **Medium**

Portfolio company logos render without explicit dimensions, contributing to Cumulative Layout Shift.

Recommendation: Add explicit dimensions to every or migrate to next/image.

Third-party hotlinked badge images on homepage/portfolio **Low**

Badges from five external directory sites add extra DNS/TLS round-trips on key pages.

Recommendation: Self-host or remove if these directory badges provide negligible referral/SEO value.

2.7 AI Search Readiness (GEO)

Score: **76/100**

What Works

- robots.txt explicitly allows every major AI crawler (GPTBot, ChatGPT-User, OAI-SearchBot, ClaudeBot, anthropic-ai, PerplexityBot, CCBot, Applebot-Extended, cohere-ai, YouBot)
- llms.txt (91 lines) and llms-full.txt (676 lines) both resolve and are content-rich, with explicit citation-provenance guidance
- Headline statistics quoted in llms.txt verified to appear verbatim in actual on-page rendered content
- Fully static, JS-independent delivery confirmed on every page sampled (mode_used: raw, is_spa: false)
- New SpeakableSpecification schema block adds voice-AI eligibility
- FAQPage schema with 26 Q&A pairs correctly typed
- Newsletter's sourced tables and Sources: attribution line are the strongest citability asset on the site

Findings

Off-site authority signals remain the primary score cap **High**

No Wikipedia entity for Kautilya or Dev Shah, no Reddit presence, and the linked YouTube channel's activity could not be verified due to a consent-wall. YouTube mentions correlate most strongly (~0.737) with AI citation odds of any signal in the GEO framework.

Recommendation: Pursue a sustained off-site footprint: YouTube video output tied to existing deal-teardown content, Reddit participation in relevant subreddits, and eventual Wikipedia notability.

Flagship stats-heavy case study has zero in-content visuals or video **Medium**

/stories/inspire3 (99.83% match rate case study) has no charts, screenshots, or diagrams supporting its own claims despite dense quantified content.

Recommendation: Add at minimum one summary-stats visual per Stories case study; embed 1-2 relevant videos on /approach and high-traffic Stories pages.

llms-full.txt excludes Stories full text Medium

llms-full.txt's header claims complete Blog and Newsletter article coverage but excludes the richest quantified case-study detail found in Stories pages.

Recommendation: Extend llms-full.txt to include full Stories/case-study text, or add a separate /llms-stories.txt referenced from /llms.txt.

News sitemap misapplies news protocol to evergreen content Medium

news-sitemap.xml still carries 2024-dated publication_date values for evergreen Stories content, coordinated with the same finding filed as High severity in the Sitemap category.

Recommendation: Remove Stories URLs from news-sitemap.xml entirely unless pursuing actual Google News inclusion.

FAQ answers remain snippet-length only Low

FAQ answers run 26-53 words, good for Google AI Overview/voice extraction but thin for Perplexity-style multi-source synthesis, which favors longer self-contained passages.

Recommendation: Link each short FAQ answer to a 134-167 word expanded answer in corresponding blog/approach content.

3. Action Plan

Phase 1: Critical Fixes Week 1

- Fix the stuck 'Loading Sky...' homepage bug creating ~17,500px of dead scroll space
- Migrate images to next/image and re-export/recompress the 16 sitemap images over 1MB
- Remove deprecated HowTo JSON-LD from /approach and /stories/[slug] pages

Phase 2: High-Impact Improvements Weeks 2-3

- Fix apex domain redirect to permanent 301/308
- Remove news-sitemap.xml
- Add width/height to portfolio-logo images
- Fix mobile CTA visibility gap (persistent ENGAGE button)
- Ship standalone Buy-Side Due Diligence service page
- Build Off-Market Deal Sourcing and Due Diligence Checklist content pillars
- Fix search-fund content cannibalization
- Kick off off-site authority building (YouTube, Reddit)

Phase 3: Content & Authority Month 2

- Fix llms.txt FAQ count and extend llms-full.txt to cover Stories
- Add @id linkage between Organization and ProfessionalService schema, reconcile pricing drift
- Add CSP header
- Add IndexNow support
- Trim over-length meta descriptions
- Add H2 structure to /portfolio and /newsletter
- Add visuals/tables to stats-heavy Stories pages
- Add author bio blocks and diversify bylines
- Add visible phone/address contact signals
- Repair internal link graph between newsletter, blog, and stories
- Expand homepage body copy and trust bar
- Remove or self-host third-party directory badges
- Fix NewsArticle schema type on Stories pages
- Consolidate image-sitemap data source

Phase 4: Monitoring & Iteration Ongoing

- Derive lastmod from actual content dates for aggregator pages
- Remove unused priority/changefreq sitemap tags
- Fix homepage canonical trailing slash
- Remove empty servicePhone schema property
- Add expanded FAQ answer variants for Perplexity-style synthesis
- Fix or restyle the homepage 'ENTER' scroll cue
- Add sticky CTA reinforcement during homepage scroll
- Reclassify legacy personal-journal Stories as Origin Story sub-section
- Add a RelatedContent component driven by cluster tags
- Re-verify performance findings against real CrUX/PSI field data once API access is available

Data Sources & Methodology

Source	Description	Update Frequency
PageSpeed Insights API	Lighthouse lab audit (mobile emulation, Moto G Power, slow 4G)	Real-time
Chrome UX Report (CrUX)	28-day rolling field data from real Chrome users	Daily ~04:00 UTC
CrUX History API	25-week p75 trend data per metric	Weekly
Google Search Console	Search Analytics (clicks, impressions, CTR, position)	2-3 day lag
URL Inspection API	Per-URL index status, coverage state, crawl info	Real-time (2,000/day)

Report generated by Claude SEO — Google SEO Intelligence Skill — August 18, 2026
Methodology based on Google Web Vitals thresholds, Search Console documentation, and Lighthouse scoring algorithms.